

Documentation – A crowdsourcing plugin for the validation of ontologies for the Protégé Ontology Editor

The crowdsourcing/GWAP plugin for the Protégé Ontology Editor allows users to validate parts of an ontology with the help of crowdsourcing – using the uComp Human Computation (HC) services. The uComp HC services rely on so-called Games With a Purpose (GWAP) and virtual labor markets such as CrowdFlower.

GWAPs are interactive games which are run as applications on Social Web platforms (such as Facebook) or as stand-alone applications. Players solve certain problems which cannot be solved by algorithms (with acceptable accuracy), GWAP users are motivated by scoring system or other forms of rewards, whereas users on virtual labor markets are paid for their input.

To use the plugin, users select the part of the ontology to validate and define some parameters (for details see below). The plugin sends the request to the HC API (uComp API), which delegates the task to a GWAP or CrowdFlower. After enough HC users given their input on the task, the results are sent back to the plugin, which then displays them to the Protégé user. Depending on the result, the user can perform further actions like deleting negatively validated parts of the ontology.

Installing the plugin

Step 1: Installing Protégé 4.3

To run the plugin, you need a recent version of Protégé installed. The current version (May 2014) is Protégé 4.3 (build 304), the download and installation instructions are found on

http://protege.stanford.edu/download/protege/4.3/installanywhere/Web_Installers/

.

The installer is available for various platforms and also with and without the Java VM. Since the version including the Java VM is not updated with every Java Update, it is recommended to install the latest Java VM separately and then download and install the Protégé version without Java. The latest Java VM is available at the Java website <http://www.java.com/de/download/>. We recommend to install the 32-bit version of Protégé. The plugin has been tested with Protégé 4.3 and 4.2.

Step2: Installing the Crowdsourcing plugin

Every Protégé plugin consists of a single jar-File (Java archive). All required third party libraries have to be included into that jar-File, but it is possible to reference to other installed Protégé plugins since Protégé is built very modular. The Crowdsourcing plugin can be downloaded from within Protégé. The relevant menu item for this functionality is found under *File → Check for plugins*.

Step3: Adding the uComp-API settings

To use the plugin, you need to provide some settings (uComp-API key, number of needed judgements per unit, payment in cents per unit if you want to publish them onto crowdfunder) in a file named **ucomp_api_settings.txt** in the folder **.Protege**, which is automatically created by Protégé in your home-directory.

You have to provide the settings in one line in the following format:

<ucomp_api_key>,<number_of_judgements_per_unit>,<payment_in_cents>

(Example file contents of ucomp_api_settings.txt:

```
abcdefghijklmnopqrst,5,2
```

This means that per unit (evaluation task) 5 different judgments will be collected, and for each judgment 2 USD-cents are paid.

)

To get a uComp-API key, request it from the uComp team, see <http://soc.ecoresearch.net/facebook/election2008/ucomp-quiz-beta/api/v1/documentation/>.

You need to specify your CrowdFlower key which will then be used by the uComp API.

General Information

Upon start, Protégé first loads the core framework which then checks and loads all plugins. In addition to the user interface of Protégé an associated console windows opens at every start. It works as standard output and log for Protégé. This helps to see which plugins are loaded by Protégé and if there are any errors. At every start Protégé checks for updates for all installed plugins. By default they are installed into the home directory of the current user. (*For Windows users:* By giving the security principal „Everyone“ full control in the NTFS-permissions of the plugins-Folder (e.g. C:\Program Files\Protege_4.3\plugins) it's possible to allow Protégé to overwrite the original jar-File.)

Until now, five types of validation tasks have been implemented, which are described in the following.

Crowdsourcing Class Validation

The screenshot shows a dialog box titled "uComp Class Validation: Accommodation". It contains the following fields and controls:

- Concept to be validated:** A text box containing "Accommodation".
- Validate relevance for domain:** A text box containing "Tourism".
- Additional information for validators:** An empty text box.
- ☒ **Send to CrowdFlower (tick) or uComp-Quiz (don't tick)**
- ☒ **Validate subtree** ☐ **Cancel subtree**
- Results area:** A text box containing:
Validating concept "Accommodation" against domain "Tourism" with additional info ""
Results from uComp quiz: Yes: 1, No: 0, I don't know: 1
uComp Validation POSITIVE! The concept "Accommodation" is relevant for the domain "Tourism!"

Crowdsourcing Class Validation lets users validate one or more ontology classes against a domain at the same time, ie. check if a class (concept) is relevant for the domain given.

Usage:

1. Load the ontology and select the "Classes" tab. From the menu, select "Window" → "Views" → "Class Views" → "uComp Class validation" and add it to the user interface.
2. The user selects an ontology class by clicking on it in the class hierarchy (i.e. „Accommodation“).
3. Define the domain against which the class is to be validated (i.e. „Tourism“) and optionally give some additional information for the validators. This additional information helps the HC users (the validators) to better understand the task and will be displayed to them, in addition to the basic task description which is included in the Plugin.
4. The user decides if the task should be sent to CrowdFlower or the uComp-Quiz (GWAP).
5. With a click on the GO-button the validation is started. Status information is displayed in the text area below.
6. The plugin sends all the necessary data to the crowdsourcing API, which creates a new crowdsourcing task. After the task has been finished, the plugin receives the results from the uComp API and displays it to the user in the text area.

7. During the whole task the user can cancel the current validation (also with the option to cancel the whole subtree). In this case, the user interface is reseted and a cancellation request is sent to the API.

Remarks:

- During the first validation, the given domain is saved in the ontology for further validations. That means if the user decides to validate another class, the textfield for the domain is automatically populated by the plugin (of course it's editable to give the user the opportunity to change it).
- If the user wants to validate **more than a single class**, (s)he can select the checkbox for the validation of the class subtree. In this case, every subclass of the current class also gets validated (recursively!) If you want to validate the **whole ontology**, select the *uppermost class* ("*Thing*"), and select option "Validate subtree". Then, all concepts will be validated.
- The information shown in the user interface only refers to the currently selected class. Therefore, when a user has chosen to validate the whole subtree, he has to select another class in the subtree to show the status and results for that class.

Crowdsourcing Subclass Validation

The screenshot shows a window titled "uComp Subclass Validation: Safari". It contains several input fields and checkboxes. The "In the domain of:" field is set to "Tourism". The "Is class:" field is set to "Safari". The "a subclass of the superclass(es):" field is set to "Adventure, Sightseeing". The "Additional information for validators:" field is set to "123". There are three checkboxes: "Send to CrowdFlower (tick) or uComp-Quiz (don't tick)" (unchecked), "Validate subtree" (unchecked), and "Cancel subtree" (unchecked). Below these are buttons for "GO" and "CANCEL". A text area at the bottom displays validation results for "Safari" against "Adventure" and "Sightseeing". A "REMOVE ALL NEGATIVE RELATIONS" button is at the bottom.

uComp Subclass Validation: Safari

In the domain of: Tourism

Is class: Safari

a subclass of the superclass(es): Adventure, Sightseeing

Additional information for validators: 123

☐ Send to CrowdFlower (tick) or uComp-Quiz (don't tick)

☐ Validate subtree GO ☐ Cancel subtree CANCEL

Validating subclass "Safari" against superclass(es) "Adventure, Sightseeing" in the domain of "Tourism" with additional info "123"

Results from uComp quiz for superclass "Adventure": Yes: 1, No: 0, I don't know: 0
uComp Validation POSITIVE! The subclass "Safari" is relevant for the superclass "Adventure"!

Results from uComp quiz for superclass "Sightseeing": Yes: 0, No: 1, I don't know: 0
uComp Validation NEGATIVE! The subclass "Safari" is not relevant for the superclass "Sightseeing"!

REMOVE ALL NEGATIVE RELATIONS

The Crowdsourcing Subclass Validation enables a user to validate **subclassof** relations between a class and its superclasses and works similar to the GWAP Class Validation. Therefore we will mention only the differences between the two.

Usage:

1. Load the ontology and select the "Classes" tab. From the menu, select "Window" → "Views" → "Class Views" → "uComp SubClass validation" and add it to the user interface.
2. When the user selects a class from the class hierarchy, the plugin automatically looks for all of its **superclasses** and puts the labels of the class and all of its superclasses into the first two textfields.
3. Again, users can provide additional information and have the option to select the whole subtree to be validated.
4. When a class has more than one superclasses (as in the example above), a validation task for each superclass will be created.
5. All results are displayed in the text area. The example also shows what happens if the result includes a negative validation: A button is displayed, which allows the user to delete all subclassof relations validated as „not relevant“.
6. If you want to **validate all subclass relations** in the ontology, go select the uppermost class ("Thing") and select option "validate subtree".

Crowdsourcing Individual Validation

The screenshot shows a dialog box titled "uComp Individual Validation: FourSeasons". It contains four text input fields: "In the domain of:" with "Tourism", "Is individual:" with "FourSeasons", "an instance of the following class(es):" with "LuxuryHotel", and "Additional information for validators:" which is empty. Below these fields is a checkbox labeled "Send to CrowdFlower (tick) or uComp-Quiz (don't tick)". There are "GO" and "CANCEL" buttons. At the bottom, a scrollable text area displays the validation results: "Validating individual 'FourSeasons' against class 'LuxuryHotel' in the domain of 'Tourism' with additional info """, "Results from uComp quiz for type (class) relation 'LuxuryHotel': Yes: 1, No: 0, I don't know: 0", and "uComp Validation POSITIVE! The individual 'FourSeasons' is relevant for the concept 'LuxuryHotel!'".

uComp Individual Validation: FourSeasons

In the domain of: Tourism

Is individual: FourSeasons

an instance of the following class(es): LuxuryHotel

Additional information for validators:

☐ Send to CrowdFlower (tick) or uComp-Quiz (don't tick)

GO CANCEL

Validating individual "FourSeasons" against class "LuxuryHotel" in the domain of "Tourism" with additional info ""
Results from uComp quiz for type (class) relation "LuxuryHotel": Yes: 1, No: 0, I don't know: 0
uComp Validation POSITIVE! The individual "FourSeasons" is relevant for the concept "LuxuryHotel!"

The Crowdsourcing Individual Validation lets users validate the relation of an OWL individual to its corresponding OWL class, ie. if the `instanceOf` relation is valid. This task can be seen as a mixture of the two previous validation tasks when looking at the validation scenarios.

Usage:

1. When a Protégé user selects an individual, the plugin automatically puts its name and the corresponding class in the first two text fields.
2. As in Subclass Validation it is possible that an individual belongs to more than one class. In this case a validation for each of the classes will be triggered. The difference to the class validation is that no subtree validation is possible, as an individual can have no sub-individuals.
3. Again, it is possible to cancel the current validation in progress.
4. If a validation task has a negative result, the respective individual can be deleted.

Crowdsourcing Object Property Validation

The screenshot shows a dialog box titled "uComp Object Property Validation: hasAccommodation". It contains several text input fields and a checkbox. The fields are: "In the domain of:" with the value "Tourism", "Object Property to be validated:" with the value "hasAccommodation", "Is this the correct rdfs:domain?" with the value "Destination", and "Is this the correct rdfs:range?" with the value "Accommodation". There is also an empty field for "Additional information for validators:". Below these fields is a checkbox labeled "Send to CrowdFlower (tick) or uComp-Quiz (don't tick)". At the bottom of the input section are "GO" and "CANCEL" buttons. A scrollable text area at the bottom displays the following text: "Validating object property 'hasAccommodation' in the domain of 'Tourism' with additional info ''", "Validating relevance of Subject (Domain) 'Destination'", "Results from uComp quiz for relevance of Subject (Domain): Yes: 0, No: 1, I don't know: 0", "Comp Validation NEGATIVE! The Subject (Domain) 'Destination' is not relevant for the object property 'hasAccommodation'", "Validating relevance of Object (Range) 'Accommodation'", and "Results from uComp quiz for relevance of Object (Range): Yes: 1, No: 0, I don't know: 0". At the very bottom is a button labeled "REMOVE NEGATIVE DOMAIN/RANGE".

With the Crowdsourcing Object Property Validation users can validate the **domain/range** restrictions of a given object property. It differs from the previous task by having more elements and also in the way the question for the Crowdsourcing users is formulated. In contrast to the other validation tasks there are two questions per object property (to check the correctness of both the domain and the range restriction).

One must distinguish between the general domain (of knowledge) of the ontology and the **Domain restriction** of a class, which should be trivial for an ontology engineer.

Usage:

1. The Protégé user selects an object property, the plugin automatically fills out the text fields with the label of the object property and its *subject (domain) and object (range)*.
2. If the general domain is already defined for the ontology, the plugin puts it into the respective text field ("In the domain of").
3. When the validation is finished, the plugin displays the results of the two validations. Again, if there was a negative validation result, the Protégé user has the option to delete the respective relations.

Crowdsourcing Relation Validation

The screenshot shows a web-based dialog box titled "uComp Relation Label Suggestion: relation". It contains several input fields and a list of checkboxes. The "In the domain of:" field is set to "Tourism". The "Suggest an object property for:" field is set to "relation". Below this is a checkbox labeled "Suggest for all relations" which is unchecked. The "Suggest between Subject (Domain):" field is set to "Activity" and the "and Object (Range):" field is set to "Destination". Below these fields is a section titled "Choose from existing object properties:" with a list of checkboxes: "isOfferedAt", "hasAccommodation", "hasActivity", "hasContact", "hasRating", and "hasPart". All these checkboxes are checked. Below this list is a text area for "Additional information for validators:". At the bottom of the form are two buttons: "GO" and "CANCEL". Below the buttons is a text box containing the following text: "Validating relation 'relation' between subject 'Activity' and object 'Destination' against the object properties: 'isOfferedAt, hasActivity'. Results from uComp quiz for suggestions for an object property: isOfferedAt: 1, hasActivity: 0, hasContact: 0, hasRating: 0, hasPart: 0. Suggested object property from uComp quiz: isOfferedAt". Below this text box is a progress bar. At the very bottom of the dialog is a button labeled "LABEL RELATION WITH SUGGESTED OBJECT PROPERTY".

uComp Relation Label Suggestion: relation

In the domain of:

Suggest an object property for:

☐ Suggest for all relations

Suggest between Subject (Domain):

and Object (Range):

Choose from existing object properties:

☒ isOfferedAt ☒ hasAccommodation

☒ hasActivity

☒ hasContact

☒ hasRating

☒ hasPart

Additional information for validators:

☐ Send to CrowdFlower (tick) or uComp-Quiz (don't tick)

Validating relation "relation" between subject "Activity" and object "Destination" against the object properties: "isOfferedAt, hasActivity".
Results from uComp quiz for suggestions for an object property: isOfferedAt: 1, hasActivity: 0, hasContact: 0, hasRating: 0, hasPart: 0.
Suggested object property from uComp quiz: isOfferedAt

This task is relevant especially for ontologies which have been created by automated methods, such as ontologies from ontology learning systems. Often ontology learning systems can only detect an unlabeled relation between two classes. The goal of this Crowdsourcing task is to select an appropriate relation label between the two classes from a given list of labels (object properties).

Usage:

1. The plugin allows Protégé users to select one or all of the mentioned unlabeled relations. The user also selects which of the existing object properties should be sent to the HC API as candidate relation labels for the unlabeled relations.
2. The plugin then sends this information to the HC API. The HC users now have to choose which of the available object properties is most appropriate for the relation between the two objects.
3. The result is sent back to the plugin and displayed to the user, who can then accept the decision in order to replace the unlabeled relation with the chosen object property.